

Masking Algorithms Guide

How each algorithm transforms data, with which parameters, and what to expect in the output.

31

ALGORITHMS

9

CATEGORIES

62

BUILT-IN INSTANCES

Contents

Algorithms are grouped using the same categories shown in the Algorithm Tester sidebar. Every input → output pair in this guide was produced by actually running the algorithm — none of them are illustrative.

PART 1 · PERSONAL DATA

1.1	Email	email.Email
1.2	Phone	phone.Phone
1.3	Name	name.Name
1.4	Full Name	name.FullName
1.5	Financial ID BR (CPF/CNPJ)	financialId.FinancialIdBr

PART 2 · DATES

2.1	Date Shift	dateAlgorithms.DateShift
2.2	Date Shift Discrete	dateAlgorithms.DateShiftDiscrete
2.3	Date Replacement	dateAlgorithms.DateReplacement
2.4	Min/Max Date/Time	minMax.MinMaxLocalDateTime

PART 3 · NUMERIC

3.1	Numeric Mapping	characterMapping.NumericMapping
3.2	Min/Max BigDecimal	minMax.MinMaxBigDecimal
3.3	Numeric Expression	expression.NumericExpression
3.4	Repeat First Digit	repeatFirstDigit.RepeatFirstDigit

PART 4 · STRING

4.1	Character Mapping	characterMapping.CharacterMapping
4.2	Character Replacement	characterReplacement.CharacterReplacement
4.3	Segment Mapping	segmentMapping.SegmentMapping
4.4	Regex Decompose	decompose.RegexDecompose
4.5	String Algorithm Chain	stringAlgorithmChain.StringAlgorithmChain
4.6	Shuffle	shuffle.Shuffle

PART 5 · FREE TEXT

5.1	Free Text Redaction	freeTextRedaction.FreeTextRedaction
5.2	Redact	redact.Redact

PART 6 · FINANCIAL

6.1	Payment Card	characterMapping.PaymentCard
6.2	IBAN	iban.IBAN

PART 7 · LOOKUP & MAPPING

7.1	Secure Lookup	secureLookup.SecureLookup
7.2	Null-Safe Secure Lookup	nullSecureLookup.NullSecureLookup
7.3	Data Cleansing	dataCleansing.DataCleansing
7.4	Mapping	mapping.Mapping

PART 8 · MULTI-COLUMN

8.1	Multi-Column Condition	conditional.MultiColumnCondition
8.2	Multi-Column Address	address.MultiColumnAddress

PART 9 · OTHER

9.1	Check Digit	checkdigit.Checkdigit
9.2	Tokenization	tokenization.Tokenization

APPENDICES

A	Cross-cutting concepts: determinism, key, references
B	Built-in instance catalog (dlpx-core:)
C	Standalone runner limitations

Personal Data

Algorithms that understand the semantics of the field — email, phone, name, CPF/CNPJ — and generate synthetic substitutes that still look like (and validate as) real data.

1.1 Email

algorithm.plugin.email.Email **DETERMINISTIC**

Masks an email address by treating each half separately: the part before the @ (the name) and the domain after it. Each half has its own strategy, which lets you fully anonymize the user while keeping the domain — preserving the per-company distribution your reports rely on.

INPUT FORMAT

A valid email address, with @ and a domain. E.g. john.smith@company.com . Malformed input falls through to errorHandlingAction .

INPUT → OUTPUT

INPUT	OUTPUT
john.smith@company.com	→ TEQVTCLYKX6Z5EGR5LYA6E44BVU0SA7ITPUJTPGZCYDNNF5MMDYA@example.com
john.smith@company.com	→ TEQVTCLYKX6Z5EGR5LYA6E44BVU0SA7ITPUJTPGZCYDNNF5MMDYA@company.com

First row uses domainAction: REPLACEMENT ; the second uses domainAction: PRESERVE . Note the generated name is identical in both — it depends only on the input and the key.

PARAMETERS

nameAction *	How to mask the name part. UNIQUE — generates a long unique identifier per input (guarantees two different emails never collide). LOOKUP — draws a name from a lookup file. APPLY_ALGORITHM — passes the name through another algorithm. APPLY_FIRSTNAME_AND_LASTNAME_ALGORITHMS — splits the name and applies first-name and last-name algorithms separately, producing something like ana.silva@... .
domainAction *	How to handle the domain. REPLACEMENT — swaps it for domainReplacementString . APPLY_ALGORITHM — passes the domain through another algorithm. PRESERVE — keeps the original domain intact.
domainReplacementString	The substitute domain, used only when domainAction = REPLACEMENT . E.g. example.com .
errorHandlingAction	What to do when the input is not a valid email. EXCEPTION — fails with an error. CHARACTER_MAPPING — applies generic character-by-character masking. PRESERVE — returns the original value unchanged.

Choosing between them. If the goal is simply to prevent accidental email delivery, domainAction: REPLACEMENT pointing at a domain you control is safest. If analysts still need to segment by company, use PRESERVE — but remember the domain can itself be identifying in small datasets.

1.2 Phone

algorithm.plugin.phone.Phone **DETERMINISTIC**

Generates a synthetic phone number preserving the original formatting exactly: parentheses, spaces, hyphens and digit count all stay in place. Only the digits change — and not all of them: the algorithm tends to preserve prefixes that identify the format (the area code).

INPUT FORMAT

A phone number in any formatting. Punctuation is kept positionally in the output. E.g. (555) 123-4567 , 5551234567 .

INPUT → OUTPUT

INPUT		OUTPUT
(555) 123-4567	→	(555) 418-3796
5551234567	→	5554183796

PARAMETERS

isPhoneUnique *	When true, the algorithm guarantees that two different source numbers never produce the same masked number. This matters when the phone is used as a deduplication or join key. When false, collisions are possible but the generated numbers look more natural.
-----------------	--

1.3 Name

algorithm.plugin.name.Name

DETERMINISTIC

REQUIRES FILE

Replaces a simple name — just a first name, or just a surname — with another drawn from a lookup file. The draw is a hash of (value + key), so the same name always maps to the same substitute: "John" will be "Thomas" across every table in the project, which keeps name-based joins working. For compound names use **Full Name** (§1.4).

INPUT FORMAT

A single word, no separators. E.g. John , Smith , Mary .

INPUT → OUTPUT

INPUT		OUTPUT
John	→	Thomas
MARY	→	ROBERT

With `maskedValueCase: PRESERVE_INPUT` — uppercase input produced uppercase output, even though the file contains "Robert".

PARAMETERS

LookupFile *	A text file with one name per line. It is the source of substitutes. More lines means a lower chance that two different names receive the same substitute. It can hold any list — first names, surnames, male or female names separately.
maskedValueCase	Capitalization of the result. PRESERVE_LOOKUP_FILE — exactly as it appears in the file. PRESERVE_INPUT — copies the input's case pattern. ALL_LOWER / ALL_UPPER — forces lower or upper case.
particlesToRemoveFile	A file of particles (one per line) stripped from the input <i>before</i> hashing. This makes "van Dyke" and "Dyke" yield the same substitute instead of being treated as distinct names.
particlesToPreserveFile	A file of particles that must reappear in the result. The generated substitute gets back the particles the original input had.
maxLengthOfMaskedName	Maximum length in characters. Names in the file exceeding the limit are excluded from the draw. Use this when the target column has a size constraint and you don't want truncation.
maxNumberNames	Maximum names returned in multi-value mode. For a simple single column, leave it empty.
filterAccent	When true, strips accents before hashing — making "José" and "Jose" produce the same substitute. Recommended for Brazilian datasets where accentuation is inconsistent.
inputCaseSensitive	When true, "John" and "JOHN" are distinct inputs and may receive different substitutes. The default (false) ignores case when comparing.

1.4 Full Name

algorithm.plugin.name.FullName

DETERMINISTIC

Splits a full name into first name(s) and surname, then applies a different algorithm to each part. It is an orchestrator: it masks nothing itself, it only decides where the given name ends and the surname begins, delegating each piece to the algorithm you point it at.

INPUT FORMAT

A full name separated by spaces. E.g. John Smith , Mary Anne Johnson .

INPUT → OUTPUT

INPUT

John Michael Smith

→

OUTPUT

Pink Deepak Luella

PARAMETERS

lastNameAtTheEnd

Where the surname sits. **true** — it is the last word ("John Michael Smith" → surname "Smith"). **false** — it is the first (for datasets storing "Smith, John").

ifSingleWordConsiderAsLastName

Tie-breaker for single-word input. **true** — treat it as a surname. **false** — treat it as a first name. This decides which of the two algorithms gets called.

firstNameAlgorithmRef

Algorithm applied to the first name(s). E.g. `dlpx-core:FirstName` . If omitted, first names are left *unmasked*.

lastNameAlgorithmRef

Algorithm applied to the surname. E.g. `dlpx-core:LastName` . If omitted, the surname is left unmasked.

lastNameSeparators

Additional separators beyond the space. E.g. `["-"]` makes "Mary-Smith" read as name "Mary" + surname "Smith".

maxLengthOfMaskedName

Character limit for the masked full name, to respect the width of the target column.

maxNumberFirstNames

How many first names to mask. With `1` , "John Peter Charles Smith" has only "John" masked — "Peter Charles" passes through untouched.

Mind the defaults. Both `...AlgorithmRef` parameters are optional, and when left empty the corresponding part is **not masked**. It is easy to configure only the surname and let first names leak into production — always check both.

1.5 Financial ID BR (CPF/CNPJ)

algorithm.plugin.financialId.FinancialIdBr

DETERMINISTIC

Masks Brazilian CPF and CNPJ numbers, generating synthetic values that remain **mathematically valid** — the check digits are recalculated. The type is detected automatically from the digit count: 11 is a CPF, 14 is a CNPJ. Input formatting (dots, slash, hyphen) is preserved in the output.

INPUT FORMAT

A CPF or CNPJ, formatted or not. E.g. 123.456.789-09 , 12345678909 , 11.222.333/0001-81 , 11222333000181 .

INPUT → OUTPUT

INPUT	OUTPUT
123.456.789-09	→ 976.201.860-50
11.222.333/0001-81	→ 50.009.982/0001-50
11222333000181	→ 50009982000150
12345678000190	→ Error: invalid check digits (maskInvalidInput = false)

Rows 2 and 3 show the same CNPJ with and without punctuation: the generated digits are identical, only the formatting differs.

PARAMETERS

maskInvalidInput	What to do when the incoming CPF/CNPJ fails check-digit validation. false (default) — raises an error and stops. true — masks it anyway, ignoring the invalidity. Legacy datasets commonly contain invalid documents; if that is your case, enable this or configure fallbackAlgorithm .
cmNumericRef	Algorithm used internally to shuffle the digits before the check digit is recalculated. If empty, uses the built-in numeric Character Mapping. Supply a catalog instance, e.g. dlp-core:CM Digits .
fallbackAlgorithm	Algorithm invoked when the value is recognized as neither CPF nor CNPJ — blank field, free text, unexpected length. Without it, those cases become errors.

Dates

Four distinct strategies for dates: shift while preserving intervals, shift by discrete values, replace outright, or merely clamp to a range.

2.1 Date Shift

algorithm.plugin.dateAlgorithms.DateShift **DETERMINISTIC**

Shifts the date by an amount drawn from `[minRange, maxRange]`. The draw is deterministic by key, so the same date always gets the same shift — which **preserves chronological order** across records. It is the algorithm of choice when analyses depend on intervals ("days between order and delivery").

INPUT FORMAT

ISO 8601: `yyyy-MM-dd` or `yyyy-MM-ddTHH:mm:ss`. The output always includes a time component.

INPUT → OUTPUT

INPUT	OUTPUT
2024-03-15	→ 2024-10-20T00:00
2024-03-15T14:30:00	→ 2024-03-26T14:30
2024-01-31	→ 2023-04-30T00:00

The first two use `±365 DAYS`; the third uses `±12 MONTHS`. Note the original time-of-day is always preserved.

PARAMETERS

minRange	Minimum shift, in units of <code>unit</code> . Negative allows moving backwards. E.g. <code>-365</code> with <code>DAYS</code> allows going back a year.
maxRange	Maximum shift. Positive allows moving forward. A narrow range better preserves data plausibility; a wide range increases protection.
unit	Unit of the shift: <code>SECONDS</code> , <code>MINUTES</code> , <code>HOURS</code> , <code>DAYS</code> , <code>MONTHS</code> , <code>YEARS</code> .
roll	Only relevant with <code>MONTHS</code> or <code>YEARS</code> , when the result would land on a non-existent date (January 31 + 1 month = February 31). true — shrinks to the last valid day of the month (Feb 28). false — overflows into the next month (Mar 3). When unsure, use <code>false</code> .

On interval preservation. Because the shift depends on the source date, two records with different dates get different shifts — the distance between them is **not** preserved exactly. If keeping distances intact is a requirement, the shift must be fixed per entity, which is what `dlpx-core:Date Shift Fixed` provides.

2.2 Date Shift Discrete

algorithm.plugin.dateAlgorithms.DateShiftDiscrete **DETERMINISTIC**

A variant of Date Shift where the shift comes not from a continuous range but from a **finite set of allowed values** defined in a configuration file. Useful when the business requires standardized shifts — for instance, only multiples of 7 days, to preserve the day of the week.

INPUT FORMAT

ISO 8601: `yyyy-MM-dd` or `yyyy-MM-ddTHH:mm:ss`.

INPUT → OUTPUT

INPUT	OUTPUT
2024-03-15	→ 2024-03-29T00:00

PARAMETERS

No parameters exposed in the schema — the set of shifts comes from the plugin's built-in configuration.

2.3 Date Replacement

algorithm.plugin.dateAlgorithms.DateReplacement **DETERMINISTIC**

Discards the original date and draws an **absolute** date within `[minDate, maxDate]`. Unlike Date Shift, there is no relationship between input and output beyond determinism — the original chronological order is destroyed. Use it when the date itself is sensitive and no temporal analysis depends on it.

INPUT FORMAT

ISO 8601 with time: `yyyy-MM-ddTHH:mm:ss`. E.g. `1985-07-20T00:00:00`.

INPUT → OUTPUT

INPUT	OUTPUT
1985-07-20T00:00:00	→ 1988-08-13T00:00

PARAMETERS

minDate	Earliest possible output date, formatted <code>yyyy-MM-ddTHH:mm:ss</code> . Sets the floor of the draw window.
maxDate	Latest possible date, same format. Pick a window plausible for the domain — birth dates in 1970–2000, for example.
unit	Granularity of the draw: <code>SECONDS</code> , <code>MINUTES</code> , <code>HOURS</code> or <code>DAYS</code> . With <code>DAYS</code> , the time component is zeroed.

2.4 Min/Max Date/Time

algorithm.plugin.minMax.MinMaxLocalDateTime

Not really masking, but **range clamping**: dates inside the range pass through untouched; dates outside are pulled to the nearest bound. The typical use is suppressing outliers that identify individuals — the centenarian customer, the 1900 sentinel date.

INPUT FORMAT

ISO 8601. E.g. 2024-03-15T14:30:00 .

INPUT → OUTPUT

INPUT		OUTPUT
1985-06-15T08:00:00	→	1985-06-15T08:00
1950-01-01T00:00:00	→	1970-01-01T00:00
2024-05-10T12:00:00	→	2000-12-31T23:59:59

Range 1970-01-01 to 2000-12-31T23:59:59 . The first date is inside and passes unchanged; the second is below the floor; the third is above the ceiling.

PARAMETERS

minDate *	Floor of the range. Any earlier date is replaced by this value.
maxDate *	Ceiling of the range. Any later date is replaced by this value.
nonConformingDataDefaultValue	Value returned when the input is not a valid date/time. Left empty, the algorithm raises a non-conforming data error and stops the job.

This does not anonymize. Values inside the range come out **identical** to the input. Min/Max is a tool for suppressing outliers, not for masking — never use it alone on a sensitive column.

Numeric

Transformations over numeric values: deterministic mapping into a range, range clamping, arbitrary Java expressions, and one fixed digit-repetition rule.

3.1 Numeric Mapping

algorithm.plugin.characterMapping.NumericMapping

DETERMINISTIC

Maps a number to another one within `[minValue, maxValue]` , deterministically. Each source value consistently receives the same destination, which preserves joins and distinct counts — but neither order nor magnitude.

INPUT FORMAT

A positive integer, no punctuation or decimals. E.g. `12345` .

INPUT → OUTPUT

INPUT		OUTPUT
12345	→	67253
98765	→	53534
11111	→	50708

PARAMETERS

minValue	Smallest value the output may take. Choose a floor that respects the column's semantics (e.g. <code>10000</code> to always guarantee 5 digits).
maxValue	Largest output value. The width of the range determines collision probability: narrow ranges make distinct values collide onto the same destination.

3.2 Min/Max BigDecimal

`algorithm.plugin.minMax.MinMaxBigDecimal`

The numeric counterpart of Min/Max Date/Time: clamps the value to a range, letting anything already inside pass through untouched. It serves to flatten outliers — very high salaries, extreme negative balances — that on their own identify a person.

INPUT FORMAT

A decimal or integer number. E.g. `5000.50` , `3.14` .

INPUT → OUTPUT

INPUT		OUTPUT
5000.50	→	5000.50
250	→	1000
15000	→	9999

Range `1000` – `9999` . Note the decimal scale of the in-range value is preserved.

PARAMETERS

<code>minValue *</code>	Lower bound. Smaller values are raised to this number.
<code>maxValue *</code>	Upper bound. Larger values are lowered to this number.
<code>nonConformingDataDefaultValue</code>	Value returned when the input is not numeric. Empty makes the algorithm raise a non-conforming data error.

This does not anonymize. Just like the date Min/Max, values inside the range come out identical. Combine it with another algorithm if the column is sensitive.

3.3 Numeric Expression

`algorithm.plugin.expression.NumericExpression`

The most open-ended algorithm in the set: it evaluates a **one-line Java expression** over the input value. The expression receives two variables — `input` (the value as a `BigDecimal`) and `seed` (a `long` derived from the key, for deterministic masking) — plus any constants you declare.

INPUT FORMAT

An integer or decimal number. E.g. `1000`, `3.14`.

INPUT → OUTPUT

INPUT · EXPRESSION	OUTPUT
<code>1234.56</code> <code>input.setScale(0, HALF_UP)</code>	→ <code>1235</code>
<code>99.4</code> <code>input.setScale(0, HALF_UP)</code>	→ <code>99</code>
<code>1234.56</code> <code>new BigDecimal(seed % 9000 + 1000)</code>	→ <code>7648</code>
<code>1000</code> <code>input.multiply(new BigDecimal(rate))</code>	→ <code>800.0000000000000444089209850062616169452667236328125000</code>

PARAMETERS

expression *	A Java expression returning <code>BigDecimal</code> . Available variables: <code>input</code> and <code>seed</code> . Classes need fully qualified names — <code>new java.math.BigDecimal(...)</code> , <code>java.math.RoundingMode.HALF_UP</code> .
inputType	How to interpret the input before exposing it as <code>input</code> : <code>DOUBLE</code> , <code>LONG</code> or <code>BIG_DECIMAL</code> (default).
constants	List of named constants available in the expression. Each has a <code>name</code> and a <code>value</code> (always a string). E.g. <code>name="rate"</code> , <code>value="0.8"</code> , used as <code>new java.math.BigDecimal(rate)</code> .
nonConformingDataDefaultValue	Value returned when the input is not numeric. Empty raises an error.

Watch out for floating point. The fourth row above returned `800.0000000000000444...` instead of `800`: the constant `"0.8"` went through a `double` before becoming a `BigDecimal`. For exact arithmetic, use the string constructor — `new java.math.BigDecimal("0.8")` — and finish with `.setScale(2, java.math.RoundingMode.HALF_UP)`.

3.4 Repeat First Digit

```
algorithm.plugin.repeatFirstDigit.RepeatFirstDigit
```

A fixed rule with no configuration: it takes the **last 4 digits** of the number and replaces them all with the first of those digits, repeated four times. The rest of the number stays intact. This is weak masking, suitable for low-sensitivity fields where you only want to break the exact value.

INPUT FORMAT

A numeric string of exactly **4, 9, 10 or 14 digits**, without punctuation. Other lengths are rejected.

INPUT → OUTPUT

INPUT		OUTPUT
1234	→	1111
123456789	→	123456666

In the second case the first 5 digits are preserved and the last 4 (6789) become 6666 .

PARAMETERS

None. The behavior is fixed.

String

The general-purpose algorithms for structured text. These are the plugin's most-used tools — and its most configurable.

4.1 Character Mapping

algorithm.plugin.characterMapping.CharacterMapping **DETERMINISTIC**

Swaps each character for another **from the same group**. You define the groups (lowercase letters, uppercase, digits...) and the algorithm guarantees a letter becomes a letter and a digit becomes a digit, preserving length and the "visual shape" of the data. Characters belonging to no group — spaces, accents, punctuation — pass through untouched.

INPUT FORMAT

Any string. E.g. Renée Smith , ABC123 .

INPUT → OUTPUT

INPUT	OUTPUT
Renée Smith123	→ Kyxéf Zawtv529
ABC123	→ YHE638
Mary	→ Rtut
Renée Smith123 (different key)	→ Qvtéi Dfnju271

The é and the space survived — neither belongs to a configured group. The last row shows the effect of the key: same input, different key, different result.

PARAMETERS

characterGroups	A list of strings; each string is a set of characters interchangeable with one another. E.g. ["abcdefghijklmnopqrstuvwxyz", "0123456789"] creates two isolated groups — letters only become letters, digits only digits. Characters absent from every group are never masked.
caseSensitive	When true, uppercase and lowercase are treated as separate groups, preserving the original capitalization. This requires you to declare both sets separately.
minMaskedPositions	Minimum number of positions that must actually change. Prevents the algorithm from returning a nearly identical value by coincidence of the draw.
preserveRanges	List of stretches to keep intact, each with start , length and direction . Use it to preserve meaningful prefixes or suffixes — branch code, control digit.
preserveLeadingZeros	When true, leading zeros are kept. Essential for numeric codes stored as text, where 007 and 7 are not equivalent.

Why accents don't change. If accented characters need masking, include them explicitly in a group: "aáâãäåæéêëíîïóôõöú". Otherwise they remain visible in the masked data and can help re-identify the record.

4.2 Character Replacement

`algorithm.plugin.characterReplacement.CharacterReplacement`

Applies character-by-character substitution rules: each rule declares *which* characters to detect and *what* to replace them with. Unlike Character Mapping, the substitution here is fixed (not drawn) — every filtered character becomes the same symbol. It can optionally run after another algorithm, via `stringAlgorithm`.

INPUT FORMAT

Any string. E.g. `Some example text`.

INPUT → OUTPUT

INPUT	OUTPUT
Some example text	→ R*v* *m*gyy* c*nz
Mary Johnson	→ Q*lv T*vb**h
Renée François (filterAccents: INPUT)	→ **n** Fp*ls**n

Rule: vowels → `*`. The consonants changed too because the algorithm applies its own masking before the rules. In the third row, `filterAccents: INPUT` converted `é` and `ç` to ASCII before processing.

PARAMETERS

stringAlgorithm	Algorithm applied to the input <i>before</i> the substitution rules. Its result is what passes through <code>replacementRules</code> . This lets you compose: mask with a strong algorithm, then normalize problematic characters.
replacementRules	List of rules applied in sequence. Each rule is a pair (set of characters to detect, replacement character).
↳ filteredCharacters	The characters this rule detects. Simple form, used when input and output share the same set.
↳ replacementCharacter	The single character that replaces each occurrence found.
↳ inputFilteredCharacters	Advanced form: characters detected in the input, when you want different sets for reading and writing.
↳ inputReplacementCharacter	Substitute applied to matches of <code>inputFilteredCharacters</code> .
↳ outputFilteredCharacters	Characters that, should they appear in the <i>output</i> of <code>stringAlgorithm</code> , must also be replaced.
↳ outputReplacementCharacter	Substitute applied to matches of <code>outputFilteredCharacters</code> .
requireInputChange	When true, requires at least one substitution to have happened — if no rule changed anything, the algorithm raises an error. This guards against configurations that silently mask nothing.
filterAccents	When to strip accents. INPUT — before applying the rules. OUTPUT — on the final result. BOTH — at both stages. NONE — never (default).

4.3 Segment Mapping

algorithm.plugin.segmentMapping.SegmentMapping

DETERMINISTIC

Splits the string into fixed-length segments and handles each one independently — mask, preserve, or replace with a constant. This is the right algorithm for rigidly formatted documents where different parts carry different meaning: CPF, CNPJ, postal code, barcode, account number with an embedded branch code.

INPUT FORMAT

A fixed-format string. With `autoIgnoreCharacters` enabled, separators (dots, hyphens, slashes) are detected and repositioned automatically. E.g. `123.456.789-09`.

INPUT → OUTPUT

INPUT	OUTPUT
123.456.789-09	→ 991.440.480-70
987.654.321-00	→ 684.669.226-24

Segments of 3-3-3-2 digits, all `MASK_NUMERIC`, with separators preserved in their original positions.

PARAMETERS

<code>segments *</code>	An ordered list of segments, in the sequence they appear in the value (ignoring ignored characters). The sum of the lengths must match the size of the data.
<code>length *</code>	How many characters this segment occupies. Ignored characters do not count.
<code>segmentType *</code>	MASK_NUMERIC — replaces with digits. MASK_ALPHANUMERIC — replaces with alphanumeric characters. PRESERVE — keeps the original. CONSTANT — replaces with the fixed value from <code>maskValues</code> .
<code>inputValues</code>	Set of characters accepted in this segment on input, as a continuous string. E.g. <code>"0123456789"</code> . Empty accepts any character.
<code>maskValues</code>	Pool of characters used for substitution. E.g. <code>"123456789"</code> (no zero) prevents the segment from starting with <code>0</code> . Empty uses the type's default pool.
<code>ignoreCharacters</code>	ASCII codes of the separators to skip when counting positions — <code>46</code> dot, <code>45</code> hyphen, <code>47</code> slash. They return to the output in their original positions.
<code>autoIgnoreCharacters</code>	When true, automatically detects all non-alphanumeric characters and ignores them. Removes the need to list ASCII codes by hand — recommended for CPF, CNPJ and postal codes.
<code>allowShortSegments</code>	When true, accepts input shorter than the sum of the segments, processing whatever is there. The default (false) raises an error.
<code>processPreserveBeforeIgnore</code>	Order of operations between removing ignored characters and computing PRESERVE segments. Enable it when a separator falls inside a preserved segment and positions come out shifted.

4.4 Regex Decompose

algorithm.plugin.decompose.RegexDecompose

The most flexible algorithm for structured text. You define regex patterns; the first one matching the **entire** string is applied, and each capture group (parentheses) receives an independent action — preserve, drop, redact, or pass through another algorithm. Text *between* groups is preserved automatically.

INPUT FORMAT

Any string. Patterns are tested in declaration order until one matches. If none match, `fallbackAction` applies.

INPUT → OUTPUT

INPUT · REGEX	OUTPUT
user@company.com ([^@]+) (@[^@]+) → REDACT, PRESERVE	→ ***@company.com
000001999191111 (\d{6}) (\d{9}) → PRESERVE, REDACT "X"	→ 000001XXXXXXXXX
"000001999191111" same pattern, trimInput: true	→ "000001XXXXXXXXX"
123456789 (\d{3}) (\d+) → PRESERVE, TRUNCATE	→ 123
ABC no fallbackAction	→ Error: No pattern matched and no fallbackAction is defined
ABC fallback REDACT "[INVALID]"	→ [INVALID]

PARAMETERS

maskPatterns	List of patterns tested in order . The first whose regex matches the entire string wins — the rest are ignored. This ordering is what enables conditional routing (see the tip below).
↳ regex	A regular expression that must match the entire string (implicitly anchored). Capture groups delimit the parts that receive actions.
↳ actions	One action per capture group, in the same order as the parentheses. Two groups require exactly two actions.
↳ actions.type	PRESERVE — keeps the stretch. TRUNCATE — removes it (becomes an empty string). REDACT — replaces it with fixed text or a character. APPLY_ALGORITHM — passes it through another algorithm.
↳ actions.redactString	Fixed text replacing the <i>whole</i> group, regardless of original length. E.g. "***". Only with REDACT .
↳ actions.redactCharacter	A single character replacing <i>each</i> character of the group, preserving length. E.g. "X" turns 9 digits into XXXXXXXXX. Mutually exclusive with <code>redactString</code> .
↳ actions.algorithm	Algorithm instance applied to the group when the type is APPLY_ALGORITHM . Accepts built-in catalog names (appendix B) or tests saved in the Tester.
fallbackAction	Action applied to the whole value when no pattern matches. It has the same four types, with its own <code>redactString</code> , <code>redactCharacter</code> and <code>algorithm</code> .

trimInput	Strips whitespace from both ends <i>before</i> attempting to match. The padding is restored in the output (see the third row of the table). Essential for CHAR(n) columns, which arrive right-padded with spaces.
requireMask	When enabled, raises an error if no pattern matches, instead of using the fallback. Use it when out-of-format values must be rejected rather than passed along.
maxInputLength	Maximum accepted input length. Longer values become errors. Empty imposes no limit.

Two common errors.

- ① `NullPointerException ... action.type is null` — when adding an action through the form, the **Action type** field starts empty; select it explicitly.
- ② Fallback action may not be PRESERVE when `requireMask` is set — `requireMask` defaults to **true**; to use `fallbackAction: PRESERVE`, set `requireMask: false` explicitly.

Conditional routing by content. Because patterns are tested in order, you can route a value to different algorithms depending on what it contains. To apply one algorithm when there is an @ and another when there isn't:

```
{
  "maskPatterns": [
    {
      "regex": "(.+@.+)",
      "actions": [
        {
          "type": "APPLY_ALGORITHM",
          "algorithm": {
            "name": "dlpx-core:EmailSL"
          }
        }
      ]
    },
    {
      "regex": "(.+)",
      "actions": [
        {
          "type": "APPLY_ALGORITHM",
          "algorithm": {
            "name": "dlpx-core:FirstName"
          }
        }
      ]
    }
  ]
}
```

The first pattern only matches values containing @; everything else falls to the second, which acts as a catch-all.

4.5 String Algorithm Chain

```
algorithm.plugin.stringAlgorithmChain.StringAlgorithmChain
```

Chains algorithms in sequence: the output of one becomes the input of the next. It requires at least two. Use it to compose transformations no single algorithm offers — mask then normalize, or apply two layers of substitution.

INPUT FORMAT

A string compatible with the *first* algorithm in the chain. The rest receive whatever the previous one produced.

INPUT → OUTPUT

INPUT

John Smith

→

OUTPUT

Salomon

Chain `FirstName` → `LastName` : the first algorithm produced a name, and the second treated that result as if it were a surname.

PARAMETERS

`algorithmReferences`

An ordered list of instance references, **minimum 2**. E.g. `[{"name": "dlpx-core:FirstName"}, {"name": "dlpx-core:LastName"}]`. Names come from the built-in catalog (appendix B) or from tests saved in the Tester.

Order matters, and so does the fit. Each algorithm must accept the format the previous one produced. Chaining a date algorithm after a name algorithm, for instance, fails — the second receives text it cannot parse.

4.6 Shuffle

algorithm.plugin.shuffle.Shuffle

NON-DETERMINISTIC

BATCH MODE

Redistributes values **across** records: each row receives a value that belonged to another. No value is invented — the set stays identical, only the associations change. This perfectly preserves the column's statistical distribution, making it ideal for columns used in aggregations.

INPUT FORMAT

Multiple values at once (batch mode). Each value must be at least `minimumShuffleSize` characters long.

INPUT → OUTPUT

INPUT BATCH		OUTPUT BATCH
Mary Johnson	→	Peter Miller
John Davis	→	Anna Smith
Peter Miller	→	Charles Wilson
Anna Smith	→	Mary Johnson
Charles Wilson	→	John Davis

PARAMETERS

`minimumShuffleSize`

Minimum length for a value to take part in the shuffle (default `3`). Shorter values are excluded and pass through unchanged.

Non-deterministic — by design. The permutation changes on every run. This is a security decision: a reproducible permutation would let someone re-run the shuffle and undo the anonymization. The practical consequence is that Shuffle **does not preserve referential integrity** across tables or across runs.

Free Text

For free-text fields — notes, reports, comments — where the sensitive data is embedded in prose that must stay readable.

5.1 Free Text Redaction

```
algorithm.plugin.freeTextRedaction.FreeTextRedaction
```

Scans free text for sensitive entities — by regular expression and/or by a list of terms in a file — and replaces only the matched stretches. Everything else is preserved, keeping the document readable and useful for analysis.

INPUT FORMAT

Free text of any length. E.g. `Contact John by email at john@company.com or on 555-0100.`

INPUT → OUTPUT

INPUT	OUTPUT
Contact John by email at john.smith@company.com.	→ Contact John by email at [REDACTED].

PARAMETERS

regularExpressions	A list of <code>{patternString}</code> objects with the regexes identifying stretches to redact. Each pattern is applied across the whole text, to every occurrence.
isDenyList	true — deny list: redacts what <i>matches</i> the patterns. false — allow list: redacts everything that does <i>not</i> match. The second option is safer for unpredictable text, but usually destroys readability.
regExRedactValue	Text replacing the stretches captured by the regular expressions. E.g. <code>[REDACTED]</code> .
lookupFile	A file of literal terms, one per line, searched for in the text. Useful for proper nouns and terms that regex describes poorly.
lookupFileRedactValue	Text replacing occurrences found via the file. It is independent of <code>regExRedactValue</code> , letting you visually mark the origin of each redaction.

5.2 Redact

`algorithm.plugin.redact.Redact`

A simpler, more direct take on redaction: a map of `regex` → `replacement_text` . Each map key is a pattern, and the corresponding value is what takes its place. With no regex configured, it returns the input value with no masking at all.

INPUT FORMAT

Any string. Stretches matching the regexes are replaced.

INPUT → OUTPUT

INPUT

Contact John by email at john.smith@company.com.

OUTPUT

→

Contact John by email at [EMAIL].

PARAMETERS

`regexRedact`

A map of `{regex → replacement_text}` . It allows different markers per data type — `[EMAIL]` , `[CPF]` , `[PHONE]` — in a single configuration.

An empty configuration masks nothing. With no regex in `regexRedact` , the algorithm returns the input untouched and issues **no** warning. Always confirm the map was filled in.

Financial

Financial instruments with verifiable structure: cards with a Luhn check digit and IBANs with a country checksum.

6.1 Payment Card

```
algorithm.plugin.characterMapping.PaymentCard
```

DETERMINISTIC

Masks card numbers while preserving the **BIN** (the leading digits identifying the network and issuer) and, optionally, the trailing digits used in reconciliation. The generated number keeps a **valid Luhn check digit**, so it still passes form and system validation.

INPUT FORMAT

A card number with or without spaces/hyphens. E.g. 4111 1111 1111 1111 .

INPUT → OUTPUT

INPUT	OUTPUT
4111111111111111	→ 4111533479707760
5500000000000004	→ 5500138521866444
4111111111111111 (preserve: 0)	→ 8036550265772777

In the first two, the BIN (4111 , 5500) was preserved automatically. In the third, with preserve: 0 , even the BIN changed — the card network is no longer identifiable.

PARAMETERS

preserve	How many digits from the end of the card to preserve. Typically 4 , to keep the last four that appear on statements and confirmation screens.
minMaskedPositions	Minimum number of positions that must actually change. Ensures the masked card does not end up too similar to the original.

Preserving the last 4 has a cost. Last-four plus BIN plus transaction date is often enough to re-identify a card in small datasets. Use preserve: 4 only when the business process genuinely depends on it.

6.2 IBAN

algorithm.plugin.iban.IBAN

DETERMINISTIC

Masks an IBAN while preserving the country code and recalculating the check digits, so the result remains a structurally valid IBAN. You control how many characters of the national part get shuffled.

INPUT FORMAT

An IBAN in standard format, without spaces. E.g. `GB29NWBK60161331926819`.

INPUT → OUTPUT

INPUT		OUTPUT
GB29NWBK60161331926819 (mask 20)	→	GB54EPWF99117911522013
GB29NWBK60161331926819 (mask 8)	→	GB26NWBK60161384059911

With `numCharsToMask: 8` the bank code (`NWBK`) and branch survive; with `20` , practically only the country remains. `GB` is always preserved and the checksum recalculated.

PARAMETERS

<code>validateInput</code> *	When true, validates the IBAN checksum before masking and rejects invalid input. Turn it off only if the dataset holds knowingly malformed IBANs that still need processing.
<code>numCharsToMask</code> *	How many characters to shuffle, counted from the end of the national part. Low values preserve bank and branch; high values mask everything after the country code.

Lookup & Mapping

Algorithms that substitute values by consulting an external source — a file of substitutes, a lookup table, or a persistent mapping database.

7.1 Secure Lookup

`algorithm.plugin.secureLookup.SecureLookup`

DETERMINISTIC

REQUIRES FILE

Replaces the value with a line from a lookup file, chosen by hashing (value + key). The same input always selects the same line, which gives referential consistency across tables without needing a mapping database. It is the plugin's most widely used lookup algorithm.

INPUT FORMAT

Any non-null string. The file's contents determine the output domain.

INPUT → OUTPUT

INPUT	OUTPUT
John Smith	→ Sarah
Mary Johnson	→ William
John Smith (ALL_UPPER)	→ SARAH

The third row shows that `maskedValueCase` only affects capitalization — the line selected from the file is still the same.

PARAMETERS

lookupFile *	A file with one substitute value per line. The number of lines defines output diversity: a file of 20 names will make many distinct values collide onto the same substitute.
hashMethod	How the line index is derived. SHA256 — cryptographic hash, deterministic by key (recommended). LEGACY — the old method, kept for compatibility with already-masked data. RANDOMIZE — random selection, <i>not</i> deterministic: every run changes the result.
maskedValueCase	Output capitalization. PRESERVE_LOOKUP_FILE — as in the file. PRESERVE_INPUT — copies the input's pattern. ALL_LOWER / ALL_UPPER .
inputCaseSensitive	When true, "John" and "JOHN" produce distinct hashes and may receive different substitutes. Leave it false to treat case variations as the same value.
trimWhitespaceFromInput	Strips whitespace from both ends of the input before hashing. Makes " John " and "John" land on the same substitute — important for <code>CHAR(n)</code> columns.
trimWhitespaceInLookupFile	Strips whitespace from both ends of each file line on load, preventing lines with trailing spaces from becoming distinct values.

7.2 Null-Safe Secure Lookup

`algorithm.plugin.nullSecureLookup.NullSecureLookup`

LIMITED IN THE RUNNER

A variant of Secure Lookup with explicit null handling: null input returns null instead of raising an error. It exposes no parameters — it uses the lookup file embedded in the plugin.

INPUT FORMAT

Any string, or null.

INPUT → OUTPUT

INPUT

sensitive-value

OUTPUT

→

null (in the standalone runner)

PARAMETERS

None.

In the standalone runner it always returns `null`. The algorithm depends on the Masking Engine context to resolve its embedded lookup file. A `null` result here is expected and does not indicate a configuration error.

7.3 Data Cleansing

`algorithm.plugin.dataCleansing.DataCleansing`

REQUIRES FILE

Substitutes values according to an **explicit** lookup table: you declare exactly which value becomes which. Unlike Secure Lookup, there is no hashing or drawing. In practice it is more a standardization tool than a masking one — normalizing abbreviations, unifying divergent spellings.

INPUT FORMAT

A string with a match in the file. With no match, the original value passes through **unchanged**.

INPUT → OUTPUT

INPUT		OUTPUT
CA	→	California
ny	→	New York
XX	→	XX (no match)

The second row matched despite the different case because `caseSensitive` is false.

PARAMETERS

lookupFile *	A file with one mapping per line, formatted <code>original{delimiter}substitute</code> . E.g. <code>CA,California</code> .
delimiter	Separator between original and substitute. E.g. <code>,</code> for CSV, <code>;</code> , or a tab. Omitted, it uses a tab.
caseSensitive	When true, <code>SP</code> and <code>sp</code> need separate entries in the file. False is the more practical choice in most cases.
trimWhitespace	Strips whitespace from both ends of the input and of the file lines before comparing, avoiding match failures caused by database padding.

Unmatched values pass through intact. This means an incomplete lookup table lets original data leak out **silently**. If the column is sensitive, validate the file's coverage before running in production.

7.4 Mapping

algorithm.plugin.mapping.Mapping

DOES NOT WORK IN THE RUNNER

Maintains a persistent one-to-one mapping in a database: each original value receives a unique substitute, stored and reused across every table and every future run. It is the plugin's strongest guarantee of **referential consistency** — and the only one that survives between separate jobs.

INPUT FORMAT

Any string. The algorithm queries the mapping set and, if the value is new, creates and persists a substitute.

INPUT → OUTPUT

INPUT

customer@email.com

OUTPUT

→ Error: Mapping set references are not supported

PARAMETERS

mappingSet *	Reference to the mapping set that stores the correspondences.
↳ algorithmName *	Name of the mapping set in the Masking Engine — it identifies which set to use. E.g. <code>client-name-mapping</code> .
↳ host	Host of the database holding the mapping set, when remote.
↳ port	Port of the remote mapping set database.
↳ database	Name of the remote database.
↳ schema	Schema of the remote database.
↳ isRemote	When true, uses the connection details above. False uses the Engine's local instance.
↳ mappingLookupKey	A partition key letting the same mapping set serve multiple isolated contexts — per client, per environment.
↳ propertiesRef	A properties file with the connection configuration, as an alternative to filling in host/port/database/schema individually.
ignoreCharacters	ASCII codes to strip from the input before the lookup. E.g. <code>45</code> hyphen, <code>46</code> dot — making a CPF with and without formatting point at the same mapping.

Unavailable in the Algorithm Tester. The standalone runner has no mapping database; any test returns `Mapping set references are not supported`. This algorithm can only be validated in the Masking Engine.

Multi-Column

Algorithms that receive the whole row rather than a single value — allowing the masking of one column to depend on the contents of another.

8.1 Multi-Column Condition

algorithm.plugin.conditional.MultiColumnCondition

DETERMINISTIC

MULTI-COLUMN

Chooses which algorithm to apply to each column **based on the value of a conditional column**. The classic case: if the gender column says "F", mask the name with a list of female names; if "M", with male names — keeping the record internally coherent.

INPUT FORMAT

A row with named columns. Required: `key` (the conditional column). Optional: `string1 ... string10`, `numeric1 ... numeric3`, `date1 ... date3`, `binary1 ... binary3`.

INPUT → OUTPUT

INPUT ROW

`key = "F"`
`string1 = "Mary"`

OUTPUT ROW

→ `key = "F"`
`string1 = "Karie"`

The `key` column is not masked — it only decides which condition applies. Only `string1` had an algorithm configured.

PARAMETERS

conditions	A list of conditions evaluated against the value of <code>key</code> . Each condition declares the values that activate it and which algorithms to apply to which column slots.
↳ key	List of conditional-column values that activate this condition. E.g. <code>["F", "Female"]</code> activates for either.
↳ string1...string10	Algorithm applied to column <code>stringN</code> when the condition activates. E.g. <code>{"name": "dlpx-core:FirstName"}</code> . Slots without an algorithm pass through untouched.
↳ numeric1...numeric3	Algorithm applied to the numeric columns (<code>BigDecimal</code>).
↳ date1...date3	Algorithm applied to the date columns (<code>LocalDateTime</code>).
↳ binary1...binary3	Algorithm applied to the binary columns (<code>ByteBuffer</code>).
fallbackAlgo	Algorithm applied to the string columns when <i>no</i> condition matches. Without it, the values pass through unmasked.
fallbackKey	Value assumed as the key when the <code>key</code> column is null or absent from the row.
keyCaseSensitive	When true, comparing <code>key</code> against the condition lists distinguishes upper from lower case. Default: false.
filterLength	Uses only the first (or last) <i>N</i> characters of <code>key</code> in the comparison. Useful when the conditional column has a meaningful prefix and a variable suffix.
filterDirection	Where to count <code>filterLength</code> characters from: <code>LEFT</code> (the start) or <code>RIGHT</code> (the end).

Where the names `key`, `string1` ... come from. They are **fixed slots** of the algorithm, not column names from your database. The translation between the real column (`GENDER`, `NAME`) and the slot (`key`, `string1`) is done in the Masking Engine **inventory**. In the Algorithm Tester you simulate that mapping by filling in the input table by hand.

Columns without an algorithm are not masked. If a condition declares only `string1`, every other column in the row passes through intact — even if it holds sensitive data. Configure `fallbackAlgo` or declare each slot explicitly.

8.2 Multi-Column Address

`algorithm.plugin.address.MultiColumnAddress` **REQUIRES ADDRESS FILE**

Masks address fields spread across several columns — street, number, district, city, state, postal code — generating a synthetic address that is **internally coherent**. Without it, masking each column in isolation would produce impossible combinations (a California street with a New York ZIP code).

INPUT FORMAT

The row's address columns. Requires an address lookup file in Delphix's specific format.

INPUT → OUTPUT

INPUT	OUTPUT
123 Flower Street	→ Error: null FileReference — address file not configured

PARAMETERS

Not documented in this guide — the algorithm does not initialize without the address file, which prevents inspecting its schema in the runner.

Not testable in the standalone runner. The algorithm fails during initialization without the address lookup file, which does not ship with the plugin. Configure it and validate directly in the Masking Engine.

Other

Two specialized algorithms: generic check-digit recalculation and reversible tokenization.

9.1 Check Digit

algorithm.plugin.checkdigit.Checkdigit **DETERMINISTIC**

Masks numbers carrying a check digit and **recalculates that digit** afterwards, so the result still passes validation. It is generic: any weighted-sum-plus-modulus scheme can be described through its parameters — barcode, EAN, payment slip, registration number, state tax ID.

INPUT FORMAT

A numeric string without punctuation, including the check digit at the configured position. E.g.
7891000315507 (EAN-13).

INPUT → OUTPUT

INPUT	OUTPUT
123456789012	→ 570382872686

PARAMETERS

weightList *	Weights applied to each digit in the weighted sum, one per data digit. E.g. EAN-13 uses [1,3,1,3,1,3,1,3,1,3,1,3]. The list length must equal numDigitsForCheckdigitCalculation.
modulusNumber *	The modulus divisor. The check digit is normally (modulus - remainder) % modulus. EAN-13 uses 10; CNPJ uses 11.
checkDigitIndex *	Position of the check digit in the string, counted from 0 at the left. In a 13-digit code with the check digit last, enter 12.
calculateChecksumRightToLeft *	Direction in which the weights are applied. true — right to left (payment slips, CNPJ). false — left to right (EAN, UPC).
numDigitsForCheckdigitCalculation *	How many digits enter the sum, excluding the check digit itself. EAN-13 has 12 data digits.
swapModulusForZeroRemainder	When the remainder is zero, uses modulusNumber itself as the check digit instead of 0. Some standards (CNPJ among them) follow this rule.
checksumCalculationType	STANDARD — weighted sum modulo N (default). TFN — the Australian Tax File Number variant, with its own logic.
numericAlgorithm	Algorithm that shuffles the digits which are not the check digit. Omitted, it uses the built-in numeric Character Mapping.
alphaNumericAlgorithm	Algorithm used when the input contains letters as well as digits. Without it, alphanumeric input may fail depending on characterHandling.
fallbackAlgorithm	Algorithm invoked for invalid input when invalidInputHandling = FALLBACK_MASK. It receives the whole original value.
preserveRegex	A regex identifying parts to preserve — fixed prefixes, country codes. Matching stretches stay intact and only the remainder is masked.
inputHandlingConfig	Block controlling how the input is handled before masking.
↳ characterHandling	NUMERIC_ONLY — accepts only 0–9. STANDARD — letters use their ASCII value. ASCII_VALUE_MINUS_48 — letters use (ASCII - 48), used by standards that interleave letters and digits.
↳ invalidInputHandling	ERROR — raises an exception (default). FALLBACK_MASK — delegates to fallbackAlgorithm.

↳ shortInputHandling	Input shorter than expected: FALLBACK delegates, PAD_LEFT / PAD_RIGHT pad with <code>padCharacter</code> .
↳ padCharacter	Padding character for short input. Usually <code>0</code> .
↳ trimWhitespace	Strips whitespace from both ends before processing. Necessary for padded <code>CHAR(n)</code> columns.

9.2 Tokenization

algorithm.plugin.tokenization.Tokenization **REVERSIBLE** **DETERMINISTIC**

The only **reversible** algorithm in the set. It encrypts the value with AES and returns a token; with the same key and configuration, the original value can be recovered (in the Tester, via the *Detokenize* button — `REIDENTIFY` mode). It is the choice when the data must return to its original form somewhere in the pipeline.

INPUT FORMAT

An alphanumeric string. Characters outside the tokenizable set — symbols, spaces, hyphens — trigger the `fallback`.

INPUT → OUTPUT

INPUT	OUTPUT
4111111111111111	→ ry5PjcFAfRYtmFWLGo7dCW59bs/Ulpuq
ry5PjcFAfRYtmFWLGo7dCW59bs/Ulpuq (REIDENTIFY)	→ 4111111111111111
4111-1111-1111-1111 (with hyphens)	→ idDL+Gi/KYr7yIm1pCnRjMFKhJvDYarNNyWV

The second row demonstrates exact reversal. Note the token is **longer** than the original — it carries the initialization vector.

PARAMETERS

fallback	What to do with non-tokenizable characters. NONE — raises an error. CHARACTER_MAPPING — applies generic masking to those characters and continues.
ivLength	Size of the AES initialization vector in bytes (default <code>8</code>). Larger values increase token randomness but consume more space in the result.
cmCharacterGroups	Only with <code>fallback = CHARACTER_MAPPING</code> . Defines which characters are interchangeable in the fallback. Empty uses the default groups.
cmMinMaskedPositions	Only with <code>fallback = CHARACTER_MAPPING</code> . Minimum positions the fallback must mask, avoiding tokens where almost nothing changed.

The token does not preserve the format. Despite being described as *format-preserving*, the observed result is a Base64 string longer than the input — `4111111111111111` (16 characters) became 32. Size the target column with room to spare, or the job will fail on truncation.

The key is the secret. Anyone holding the key can reverse any token. Treat it with the same rigor as a production key: out of the code, out of the repository, with controlled rotation — and remember that rotating invalidates every token already issued.

Cross-cutting concepts

Behaviors that span several algorithms and tend to cause confusion the first time you meet them.

DETERMINISM AND THE ROLE OF THE KEY

Most algorithms are **deterministic**: the same input value, with the same key, always produces the same output. That is what preserves joins — if `John` becomes `Thomas` in the customers table, it becomes `Thomas` in the orders table too.

The key is what makes the result unpredictable to anyone who lacks it. Changing it changes the entire output, as seen in §4.1. Two practical consequences: **the same key must be used across the whole project**, and **changing the key invalidates consistency with already-masked data**.

BEHAVIOR	ALGORITHMS
Deterministic by key	Character Mapping, Numeric Mapping, Segment Mapping, Payment Card, IBAN, Check Digit, Date Shift, Date Replacement, Secure Lookup, Name, Full Name, Email, Phone, Financial ID BR, Tokenization
Non-deterministic by design	Shuffle (§4.6) — a different permutation on every run; Secure Lookup with <code>hashMethod: RANDOMIZE</code>
Key-independent	Min/Max (both), Redact, Free Text Redaction, Data Cleansing, Repeat First Digit, Numeric Expression without seed

ALGORITHMS THAT MAY MASK NOTHING

Several algorithms silently return the original value under certain conditions. None of them issue a warning — verifying is on you.

ALGORITHM	WHEN DATA PASSES THROUGH INTACT
Min/Max BigDecimal · Date/Time	Whenever the value is already inside the range
Data Cleansing	Value with no match in the lookup file
Redact	<code>regexRedact</code> empty, or no pattern matching
Regex Decompose	<code>fallbackAction: PRESERVE</code> with a value matching no pattern
Full Name	<code>firstNameAlgorithmRef</code> or <code>lastNameAlgorithmRef</code> not configured
Multi-Column Condition	Column slot without an algorithm, or no condition matching and no <code>fallbackAlgo</code>
Character Mapping	Characters outside every <code>characterGroups</code> entry (accents, punctuation)

CHAR(N) COLUMNS AND INVISIBLE PADDING

`CHAR(n)` columns right-pad the value with spaces up to the declared size. An 11-digit CPF in a `CHAR(15)` column reaches the algorithm as `"12345678909 "`. Since `\d` does not match a space, regexes that work in the Tester fail in production with messages like `No pattern matched` or `did not conform to the expected format`.

The fix is to enable the algorithm's trim option: `trimInput` on Regex Decompose, `trimWhitespace` on Check Digit and Data Cleansing, `trimWhitespaceFromInput` on Secure Lookup. The padding is restored in the output.

REFERENCES BETWEEN ALGORITHMS

Many parameters ask for an "algorithm reference" — `firstNameAlgorithmRef`, `actions.algorithm`, `numericAlgorithm`, `fallbackAlgo`. The expected value is `{"name": "<instance>"}`, where the instance can be:

① a name from the built-in catalog, prefixed with `dlpx-core:` (appendix B); ② the name of a test saved in the Algorithm Tester itself; ③ in the Masking Engine, any configured algorithm instance. Non-existent names produce `Sub-algorithm not found`.

Built-in instance catalog

The 62 instances shipped with the plugin, usable in any reference parameter. Prefix them with `dlpx-core:` — for example `dlpx-core:FirstName`. Names containing spaces work fine.

ABARoutingNumber SL	IBAN	Legacy SchoolNameLookup
Age SL	Japan Corporate Number	Legacy USStateCodesLookup
BR – Financial ID	Japan Drivers License	Legacy USStatesLookup
BloodType SL	Japan My Number	Legacy UsCitiesLookup
Brazil CNPJ	JobTitle SL	Legacy UsCountiesLookup
Brazil CPF	Language SL	Legacy WebURLsLookup
CM Alpha-Numeric	LastName	MaritalStatus SL
CM Digits	Lat_Long Coordinates	MedicalGenericDrug SL
CM Numeric	Legacy AccNoLookup	NationalOrigin SL
Countries SL	Legacy AddrLine2Lookup	Null Secure Lookup
Credit Card	Legacy AddrLookup	Phone US
Date Shift Discrete	Legacy BusinessLegalEntityLookup	Phone Unique
Date Shift Fixed	Legacy CommentLookup	PoliticalOrientation SL
Date Shift Variable	Legacy DrivingLicenseNoLookup	Redact Digits–Zero
Department SL	Legacy DummyHospitalNameLookup	ReligiousOrientation SL
Email SL	Legacy EmailLookup	Repeat First Digit
Email Unique	Legacy FirstNameLookup	SexualOrientation SL
Ethnicity SL	Legacy FullNMLookup	Shuffle
FirstName	Legacy LastCommaFirstLookup	SwiftCode SL
FullName	Legacy LastNameLookup	TimeRange
Gender SL	Legacy RandomValueLookup	

The **SL suffix** stands for *Secure Lookup* — pre-configured instances with themed lookup files embedded (job titles, countries, marital status). They are the fastest way to mask categorical columns without assembling your own file.

Standalone runner limitations

The Algorithm Tester runs algorithms outside the Masking Engine. Some features depend on the Engine's context and do not work locally.

FEATURE	STATUS
FileReference	Supported. Use a <code>file:///absolute/path</code> URI. Requires <code>commons-codec</code> in <code>lib/</code> . Files are managed under <i>Settings</i> → <i>Files</i> .
MappingSetReference	Not supported. The <i>Mapping</i> algorithm (§7.4) fails with <code>Mapping set references are not supported</code> — it depends on a mapping database.
EmbeddedGeneratorReference	Not supported.
MaskValueMetadata	Always returns <code>null</code> .
Multi-Column Address (§8.2)	Does not initialize without Delphix's address lookup file.
Null-Safe Secure Lookup (§7.2)	Always returns <code>null</code> — it depends on the embedded file resolved by the Engine.

About the outputs in this guide. Every input → output pair was generated by running the algorithms in `AlgorithmRunner` against the `delphix-algorithm-plugin 2026.3.0-SNAPSHOT` plugin. Deterministic algorithms reproduce these exact values given the same key; Shuffle, by its nature, produces a different permutation on every run.